

Statistical Learning: Feature Learning

Jesse Gronsbell

Department of Statistical Sciences, University of Toronto

Recap from yesterday

- How to compute or *estimate* m and b for the linear model with LS
- More on LS and linear regression
- Logistic regression


What we'll cover today

- Feature learning
- Model evaluation

Part III: Principal Component Analysis

What happens when we have more covariates?


temp sales



$p = 1$


vs.

image
vectors



What happens when we have more covariates?

temp sales



$p = 1$

vs.

image
vectors




image type
ex. 0/1

When p gets larger we often ask, is there a more compact representation of the covariates?

Why do we have to think about this?

- In the old days, we had a **small number of covariates** to build a model
- We now have **big data sets** that require a lot of cleaning and preparation
- **Cleaner and correctly prepared data** will lead to “**good**” results
 - Good = prediction performance, interpretation, understanding, trustworthiness, etc.

Feature (or representation) learning

- A subfield of statistical learning (or ML) focusing on methods that yield **useful and often concise representations of the data**
- Can replace **feature engineering** or the manual process of identifying covariates
- Useful when:
 - We have a **lot of covariates** and want to simplify and/or remove noise
 - The data set doesn't have a **natural representation**
 - **Ex.** Yelp dataset on kaggle has 5.2 million reviews on 174,000 businesses

There are many flavors of feature learning

Supervised

Find a representation of \mathbf{X}
using \mathbf{y}

- Linear discriminant analysis ('shallow')
- Neural network ('deep')

Unsupervised

Find a representation of \mathbf{X}
without using \mathbf{y}

- Principal component analysis ('shallow')
- Autoencoder ('deep')

There are many flavors of feature learning

Supervised

Find a representation of \mathbf{X}
using \mathbf{y}

- Linear discriminant analysis ('shallow')
- Neural network ('deep')

Unsupervised


Find a representation of \mathbf{X}
without using \mathbf{y}

- Principal component analysis ('shallow')
- Autoencoder ('deep')

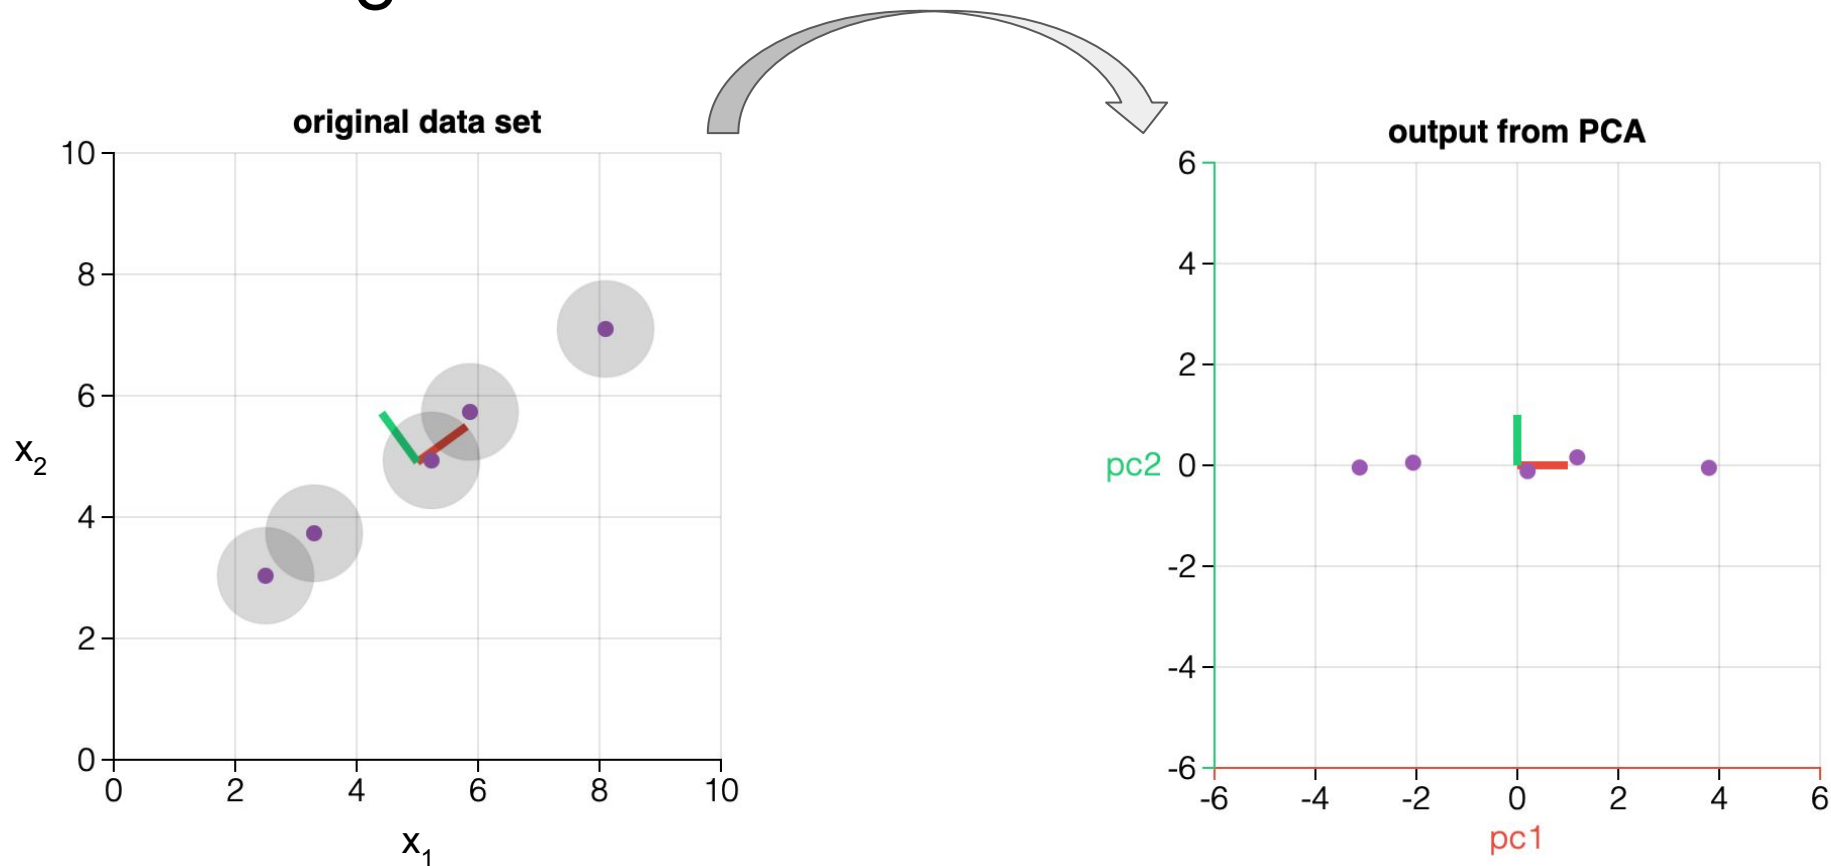
Principal component analysis (PCA)

- PCA **generates new covariates** that are weighted sums of the original covariates
- The new covariates **capture the variance** in the data and are **uncorrelated**
- Sometimes, only a **subset of the new covariates** contain all the variability in the original covariates
 - PCA yields a **more compact** representation of the original data set without losing information

Visualizing PCA



Visualizing PCA



Intuition for the math behind PCA

- We obtain the principal components with the **eigendecomposition** of the **covariance matrix** of X
- The **covariance matrix** tells us how the covariates vary with one another
- The **eigendecomposition** gives us:
 - **Eigenvectors**: the directions (or components) that capture the most variance
 - **Eigenvalues**: the magnitude of the directions
- The **eigenvectors with the largest eigenvalues** are the principal components

PCA is great - what is the catch?

- We may have a more compact data set, but our new covariates are **harder to understand**
- We pay the price in what is called **model interpretability**
- More in the deep learning lectures!

